

Harsh Amin

587-664-9195 | aminharsh317@gmail.com | harshamin.me | [in](#) hamin06 | [G](#) hamin2006

EDUCATION

University of British Columbia

BSc Co-op Computer Science · Math

GPA: 4.00

April 2027

Vancouver, BC

Awards: Trek Academic Excellence Scholarship - awarded to top 5% of faculty

Extracurricular: UBC Launch Pad Design Team, Intramural Hockey, Weight Lifting Club, Entrepreneurship Club

TECHNICAL SKILLS

Languages: Java, Python, C/C++, Racket, R, HTML/CSS, JavaScript, TypeScript, SQL, GraphQL

Frameworks: LangChain, LlamaIndex, Agents, Ollama, React, Flask, FastAPI, Node.js/Deno, Express.js, SQLAlchemy

Developer Tools: AWS, AWS CDK, REST, OpenAPI, PostgreSQL, MongoDB, Docker, Git, Linux/Unix, Jenkins

EXPERIENCE

Software Developer Intern

May 2025 – Dec 2025

AWS Cloud Innovation Centre @ UBC

Vancouver, BC

- Built open-source **Generative AI** applications **end-to-end** for **3+** sponsor teams, leveraging **AWS Bedrock, IAM, WAF, Shield, Cognito, Secrets Manager** to deliver secure and scalable **AI/ML** solutions
- Developed an **AI Agent** using Anthropic's **MCP** that maintained context across **10+** multi-turn interactions, improving dataset query accuracy by **25%+** for researchers
- Migrated from unreliable local Docker workflows to a standardized **AWS CI/CD** pipeline (**CodeBuild, CodePipeline, ECR**), cutting cross-environment inconsistencies by **33%** and reducing deployment time by **40%**
- Collaborated with a cross-functional team of **5+** developers, managers, and stakeholders, presenting technical solutions and documenting each phase for seamless deployment and knowledge transfer

Technical Lead + Developer

Sep 2024 – Present

UBC Launchpad - Project: Forum AI (LLM based student question board for UBC classes)

Vancouver, BC

- Designed the **Retrieval Augmented Generation** model and updated it to be **20%** quicker
- Co-lead a **15+** member team of developers and designers, arranging weekly syncs and tickets over 8 months
- Migrated to **FastAPI** endpoints and implemented trigger-based embedding retrieval from a **PGVector** database, resulting in a **25%** reduction in API response time
- Led the migration to the **Deno** runtime, optimizing containerized function performance and achieving a **20%+** improvement in execution speed within **Docker** environments for resource heavy functions (ex. RAG)

TECHNICAL PROJECTS

Legal Aid Tool | AWS CDK, LangChain, RAG, OpenAPI, Docker, Lambda, Bedrock, WAF, RDS

May 2025

- Developed a cloud-native application sponsored by the UBC Allard School of Law, enabling **150+** law students to streamline client interviews and reduce case-prep time by **~ 30%**
- Built a case-specific chat assistant leveraging **AWS Bedrock** and a **LangChain RAG** pipeline, scaled to process **200+** case files and **~15** hours of audio transcripts with embeddings stored in **RDS PostgreSQL**
- Deployed a production-ready **API Gateway** stack from an **OpenAPI** spec, integrating **10+** **Lambda** functions (via **Docker** images and **Lambda Layers**) with automated CI/CD integration scaling to **1K+ requests/day**

Forum AI | PostgreSQL, React.js, Next.js, Deno, Docker, AWS, SQLAlchemy

Sep 2024

- Created a Student-Instructor Q&A forum, collaborating with faculty and enhanced by Gen AI, with **200+** users
- Incorporated a **PostgreSQL** database (**Supabase**), developing migration scripts in Python with **SQLAlchemy** and **Alembic**, designing the organizational structure for project SQL queries
- Utilized **Docker** to containerize the application with support for **OCR** using Google's **Tesseract** model

Auto Pilot | Jetson Nano, YOLOv8, TensorRT, OpenCV, PyTorch, Jetson.GPIO

April 2025

- Developed an edge-based self-driving car software stack on a **Jetson Nano** to perform real-time object detection, lane following, and steering control using **deep learning** and **computer vision**.
- Trained and optimized a **YOLOv8** model using **PyTorch** for road sign and obstacle recognition; exported to **ONNX** and accelerated inference with **TensorRT** to achieve real-time performance on embedded hardware.
- Utilized **OpenCV** for **perspective transformations** to assist in path planning and steering control logic.